

Section	Comment	Source	Action	Assignee
21.14	DISCUSSION, para 1. Last sentence, RFC references have extraneous () around references.	Kim - #19 in https://www.ietf.org/mail-archive/web/dhcwg/current/msg18132.html	Ignored this one for now as we use ([<rfc>]) in other places (total of 7) when we are just adding an information reference for the reader to refer to. But maybe we should drop the parenthesis or use comma (s)? Or do we just punt this for now and let RFC-Editor deal with it?	Tomek removed extra parentheses.
Appendix B & C	Should we drop these? While I think Kim's comments and my review have fixed the issues, I just wonder if they really add value.	Kim - #26-33 in https://www.ietf.org/mail-archive/web/dhcwg/current/msg18132.html	Add text that these tables are for information purposes only and that if there is a disagreement with the text of the document, the text takes precedence.	Bernie - Done (added new text - "This table is informational and should it conflict with text earlier in this document, that text should be considered authoritative.")
18.3	<p>The sentence:</p> <p>"The server must be aware of the recommendations on packet sizes and the use of fragmentation in Section 5 of [RFC2460]."</p> <p>Suggest changing this to:</p> <p>"The server must honor the recommendations on packet sizes and the use of fragmentation in Section 5 of [RFC2460]."</p> <p>Also, this same sentence appears a second time in Section 18.3.9. Since Section 18.3 already applies to all server response message types, I believe this duplicate sentence in Section 18.3.9 is redundant and could be removed.</p>	Fred - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18147.html	<p>Bernie - I'd be more in favor of leaving this as is since I'm not sure "must" is so good (should it be MUST)?</p> <p>We should remove the redundant sentence in 18.3.9.</p>	<p>Bernie - Done (removed redundant sentence in 18.3.9.)</p> <p>Leave 18.3 text as is.</p>
8, 9.1, 9.2, 20.4.1, 21.1, 21.*	<p>Should we document the size of the message format fields? For example, msg-type (1 octet), transaction-id (3 octets), hop-count (1 octet), link-address/peer-address (16 octets)? Currently, one must view the diagram to determine this? We document the size for options (section 21), so seems odd not to have done this for the messages. But RFC3315 didn't, so perhaps OK.</p> <p>Also, noticed Type field in 20.4.1 doesn't specify size either.</p> <p>And 21.1 should specify size of option-code, option-len?</p> <p>And, turns out that 21.* sections aren't great about specifying size after all?</p>	<p>Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html</p>	Let's add text to explicitly specify the size of fields.	Bernie - Done

16.11	"the message does not include DHCP authentication ...", should we explicitly say RKAP or leave general? We might also want to just say uses RKAP or some other (yet to be developed, such as <ref to sedhcpv6>) authentication?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Bernie - I think reworking this to say "The message does not include authentication (such as RKAP, see section 20.4) or fails authentication validation."	Bernie - Done
17	Page 44 (PDF): What does this 1 paragraph section intend to say/mean? ("Client's behavior is different depending on the purpose of the configuration.") Can we drop it? Should we say something else/more?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Tomek clarified that it's about interface selection.	Done
Section 18	Page 45 (PDF): (Kim mentioned this): There's a bunch of text duplicated in this section (more than just the list). I do think "The client has many" and the list following should be removed. I also wonder if the later paragraph (The client is responsible for creating IAs ...) should be moved up as part of the Solicit exchange (but it would need to be edited/merged in to that earlier material). The duplicated instance of "A server may initiate a message exchange" should be dropped.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Tomek reviewed revised text, paragraph move up a bit, the duplicate text seems to be gone already. The list reworded slightly (to be more clear that the list is not exhaustive)	Done
Section 18.2	Page 50 (PDF) (last paragraph of section 18.2): The "and duplicate responses by servers" doesn't make any sense in this regard. Unicast can only be used for instances where a single server is being addressed so it makes no sense to say this.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Tomek ok with removing.	Bernie - Done
18.2.1	Page 51: The "after a power outage" is technically wrong as it should be "after recovery from a power outage"? We should double check if used elsewhere and correct.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html		Bernie - Changed to use "after recovery from a power outage".
18.2.4	Page 57 (last paragraph): So this says "If the terminated Renew exchange was not initiated as a result of receiving a Reconfigure message, the client begins Rebind message exchange...". But isn't this true even it was initiate by receiving a Reconfigure? The Renew is terminated at (earlist) T2 and a Rebind would start then anyway? This doesn't make sense to me. I think we should just drop the "was not initiated as a result of receiving a Reconfigure message"? Or what does happen in that case that is different?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Remove the text. Should proceed to rebind, even when caused by reconfigure.	Tomek - Done (text removed)

18.2.10	Page 63 (Top of page paragraph): "The client resends the original message using multicast". So should we qualify this to only happen if the client was using unicast (otherwise a storm could occur if "server" sends use Multicast and client retransmits using multicast)? Other option might be to say that retransmission doesn't happen until normal timer expires? Perhaps we can let client implementers chose either one depending on what's easier to implement for them?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Ignore issue.	None.
18.2.10	Page 63 (2nd to last paragraph of section): "This behavior does not apply to other IA containers, and their processing is described in detail on other parts of this document.". I'm not exactly sure where that is described. Could we add references? Or was this intended to reference other documents that might describe containers?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Replace "in other parts of the document" with "in the next section"	Tomek
18.2.11	Page 67 (top of page): "Subsequent Reconfigure messages cause the client to initiate a new exchange." Subsequent here is a bit confusing since the text preceding just said "the client MUST ignore any additional Reconfigure messages until the exchange is complete". Perhaps we should change subsequent to "After the exchange is complete, ..."?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Remove "subsequent reconfigure messages ... sentence."	Tomek - Done (text removed)
18.3	Page 68 (2nd to last paragraph on page): I think we can say "These Advertise and Reply messages MUST" (instead of just "This Reply")?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Apply change	Bernie - Done
18.3	Page 68 (top of page paragraph): I think "reply message" should be "response message" (that was used on previous page).	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html		Bernie - Done
18.3	Page 68 (1st full paragraph, end): I think the printing and new software examples probably aren't that hot? DHCPv6 has no printing options so why would printing impact client configuration. And, while software might be OK given the BOOTFILE options, but it seems unlikely. Perhaps we just say when "other configuration options are updated, perhaps because servers are moved, added, or removed" or something more general?	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Apply change	Bernie - Done

18.3.9 and 18.3.11	<p>Page 78 & 80: I suggest cleaning this up a bit as follows:</p> <ul style="list-style-type: none"> - Change 18.3.9 title to remove transmission (so just "Creation of Advertise Messages") - Change 18.3.9 to remove the last two paragraphs (perhaps replace with text to say "Transmission of the Advertise message is described in the next section.") - Change 18.3.10 title to "Transmission of Advertise and Reply Messages" - Change text in 18.3.10 from using "Reply" to "Advertise or Reply". - Add a reference to section 19.3 as that talks about how the server should respond. This is probably better reference than just to 21.10 (which could be moved to 19.3). 	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Apply change	Bernie - Done
20.2 and 20.3	Pages 85 & 86: Merge these as we could just remove the section break and the first paragraph of 20.3.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Ignore	Nothing to do.
21.6	Page 95: Drop the last sentence of first paragraph; seems broken.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Done	Bernie - Done (made changes per Kim's review (https://www.ietf.org/mail-archive/web/dhcwg/current/msg18132.html), #14)
21.24 and 21.25	Page 117 and 118: Replace "in the main body of the message to client" with "top-level option"? We have "top-level option" in terminology section.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Change	Bernie - Done. Also noticed field list for INF_MAX_RT option used SOL_MAX_RT field (should be INF_MAX_RT); fixed.
24	Page 121: Remove "IANA is requested to add a column to the DHCPv6 Option table" as next paragraph asks them to add two columns and this should be been deleted.	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Done	Bernie - Done
Appendix B	Page 136: The "Auth." Option was incorrectly marked as being used in an Information-Request, but it should have been in the Reconfigure (1 line earlier). So, this needs to be moved "up".	Bernie - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18136.html	Done	Bernie - Done

Shepherd Doc Issue	<p>>>> - I'm okay with deprecating the delayed authentication protocol, but I'm not sure if it's okay to ship the spec with no built-in authentication (the reconfigure key protocol is very limited and too weak).</p> <p>></p> <p>> BV> We are working on sedhcpv6 but this has not been so simple and quick. And, DHCPv6 is in use and deployed today without an authentication protocol. There are other actions that can be taken such as stated in section 22:</p> <p>></p> <p>> Many of these rogue server attacks can be mitigated by making use of the mechanism described in [RFC7610].</p> <p>I personally wouldn't argue that it's a blocking issue, and, to be clear, I'm not requesting sedhcpv6 be specified as *the* authentication protocol in rfc3315bis. But I'd suggest chairs to explicitly leave notes about the point for security ADs when the doc is sent to the IESG.</p>	Jinmei - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18146.html	<p>Ralph will add a note about this in the Shepherd write-up when requesting publication for this document.</p> <p>Deployment of DHCPv6 is not waiting on an authentication mechanism and is deployed without it today. The BIS document reflects what is currently in deployment.</p>	Ralph?
9.1, 19.1.1, 19.1.2, 26	<p>>>> - Section 8.1</p> <p>> [...]</p> <p>></p> <p>>>> site-scope unicast address has been deprecated. Also, the phrase</p> <p>>>> "global or ULA" is awkward as ULA is a global (scoped) address.</p> <p>></p> <p>> The 08 version states:</p> <p>></p> <p>> This is typically a globally unique address or unique local address (ULA) [RFC4193],</p> <p>></p> <p>> This is still awkward to me, since in a sense ULAs are also "globally unique" (even if the uniqueness is not 100% guaranteed). Personally I'd simply avoid mentioning ULAs unless there is a specific reason to refer to it in the context (and I don't see such a reason here), but if we want to refer to ULAs here, I'd say, e.g., "a global address (including unique local addresses)".</p> <p>></p> <p>> BV> Unless you suggest some text, we will leave it be. I don't see that this is necessarily in conflict - yes, ULA is in the "global address space". But I also think it adds value to explicitly mention ULAs here as otherwise someone could confuse the subtly that ULA is allowed.</p> <p>I've already suggested text (I admit the point may be a bit pedantic, though).</p>	Jinmei - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18146.html	<p>Bernie - I prefer leaving the text as is (global or unique local). While perhaps technically redundant, it does no harm and adds clarify. I guess we could go with "global address (including unique local address)" or something similar if others feel we should be more technically accurate?</p> <p>Tomek - I like Bernie's suggestion. I made the necessary changes.</p>	<p>Tomek - Done. Updated per Bernie's comments.</p> <p>Tomek updated - but Bernie comments: in 2 places you used "unique local address" and in one "ULA". Section 9.1 (Relay-forward Message) does define ULA (previous to your edits, it was never used – now used in 1 place). Perhaps we should cleanup as follows:</p> <ul style="list-style-type: none"> - Remove (ULA) from section 9.1 - Replace your one use of ULA with unique local address

<p>General - concern about T1/T2 handling</p>	<p>>>> - Section 17.1.10.1 > [...] >>> Specifically, the bullet description seems to lead to multiple separate state machines for different bindings. > > BV> I believe we addressed this by other text in the document - for example, see in 18.2.4 the following text: > > The server controls the time at which the client should contact the > server to extend the lifetimes on assigned leases through the T1 and > T2 parameters assigned to an IA. However, as the client Renews/ > Rebinds all IAs from the server at the same time, the client MUST > select a T1 and T2 time from all IA options, which will guarantee > that the client will send Renew/Rebind messages not later than at the > T1/T2 times associated with any of the client's bindings (earliest > T1/T2).</p> <p>Hmm...this is one of the things why I can't be confident about this spec without actually trying to implement it from the text. But I don't have enough time to check if the concern is addressed in the entire context, I'll simply give up with it.</p>	<p>Jinmei - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18146.html</p>	<p>Bernie - I think we are OK with this, but a critical review to see if it is fully handled would not hurt?</p>	<p>Ignore</p>
---	---	---	--	---------------

21.22 (IA Prefix option)	<p>>>> - Section 20.22 > [...] >>> This choice of valid/preferred lifetime in an IA prefix and the >>> (possible) relationship between them and those lifetimes advertised >>> in the delegated site do not make perfect sense to me. > [...] > > I don't think the revised version addresses the main point. The new > text simply avoids saying specifics about these lifetimes, but then > it's even more unclear about what these lifetimes are (especially > the preferred lifetime). > > BV> Unless you provide text, I think we'll leave this alone as > mentioned earlier we feel that this document is not the one that > should describe how these lifetimes are to be used.</p> <p>This is not just about wording. It's about actual protocol behavior, and I don't even know if we have common consensus on it. I also think we can't simply ignore the issue by being silence, since it could result in a bad situation like a site keeps advertising a prefix in RA as a preferred one even when it's already "deprecated" or even "invalidated" in terms of prefix delegation. I'll see if I can offer something more concrete, but I'd like to raise it as a possible blocking issue.</p>	Jinmei - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18146.html	Bernie - Updated per email discussion with Jinmei. See Row 29.	
--------------------------	--	--	--	--

21.24 (SOL_MAX_RT) and 21.25 (INF_MAX_RT)?	<p>>>> - Section 20.23</p> <p>> [...]</p> <p>>>> The expected usage is not very clear to me here. Is it intended to</p> <p>>>> be used for Advertise and apply the value to the ongoing</p> <p>>>> Solicit-Advertise exchange? Or is it mainly intended to be used for</p> <p>>>> a possible subsequent restart of a DHCPv6 session? [...]</p> <p>> BV> You original question was I think about why SOL_MAX_RT MUST be</p> <p>> in ORO. It MUST be as per RFC 7083 so that clients can get updated</p> <p>> SOL_MAX_RT values. And, yes, the idea is that the Advertise with the</p> <p>> SOL_MAX_RT value is used (even if it offers no leases) so that the</p> <p>> client applies this value for any future requests. I don't see the</p> <p>> need for this to be spelled out?</p> <p>I'd think it's worth explaining, but I'll leave it to you.</p>	Jinmei - https://www.ietf.org/mail-archive/web/dhcwg/current/msg18146.html	<p>Bernie - I'm not sure what is unclear about this. Text (for SOL_MAX_RT) is:</p> <p>If a DHCP client receives a message containing a SOL_MAX_RT option that has a valid value for SOL_MAX_RT, the client MUST set its internal SOL_MAX_RT parameter to the value contained in the SOL_MAX_RT option. This value of SOL_MAX_RT is then used by the retransmission mechanism defined in Section 15 and Section 18.2.1.</p> <p>This seems pretty clear that it applies ASAP (should be applies now and retained for as long as possible to be used in future Solicit exchanges).</p>	Tomek - agree, the text is pretty clear to me. No change needed.
6.3 and 18.2.10.1	From discussions with Jinmei and Tim Winters regarding PD and router advertisements "lifetimes". Noticed that text says IA_PD when it should have used IAPREFIX!			Bernie - Done
Appendix A	We should likely have added something about clarifying the use of preferred/valid lifetimes for IAPREFIX - Jinmei's issues above.	Bernie	Do this for an -10 version!!	